

Introduction

The **Game of Atlas** is a word-based game that can be represented as a **directed graph**, where:

- **Nodes represent countries.**
- **Edges exist if the last letter of one country matches the first letter of another.**

Using this representation, we can apply **graph theory and deep learning techniques** to **analyze, optimize, and strategize our gameplay**. This report outlines **four strategic approaches**:

1. **Out-Degree Differential Strategy** – Restricting the opponent's viable moves.
2. **Entropy Minimization and Forced Predictability** – Forcing opponents into predictable patterns.
3. **Centralized Choke-Point Strategy** – Controlling key transition nodes.
4. **Local Lookahead Strategy** – Simulating multi-step moves for long-term dominance.

Additionally, we explore **two community detection techniques**:

1. **Infomap**, an information-theoretic approach that models the game as an **information flow system**.
2. **GNN + HDBSCAN**, a deep learning-based method that learns **graph embeddings and clusters countries dynamically**.

Each strategy was **simulated in a 20-move game**, with results visualized through **performance metrics and interactive graphs**.

Task 1: Competitive Advantage Strategies

1. Out-Degree Differential Strategy

Goal: Maximize your own move options while ensuring that the opponent has significantly fewer viable responses.

Failed Sub-Strategy: Maximizing Your Own Out-Degree

Idea:

A simple approach would be to **always choose the country with the highest out-degree**, ensuring maximum follow-up options.

Why It Failed:

- **Doesn't account for opponent moves:** The opponent may also have many strong responses.
- **Leads to an evenly matched game:** No **long-term control** over the board.
- **May benefit the opponent:** High-out-degree nodes might give the opponent multiple options too.

Improvement: Out-Degree Differential Strategy

Instead of blindly maximizing out-degree, this strategy:

- **Minimizes the opponent's best possible out-degree.**
- Ensures that even if the opponent picks their best response, **we retain a move advantage.**

Simulation Results:

Both Strategic:

Player 1 Win Rate: 49.10%

Player 2 Win Rate: 50.90%

P1 Strategic vs P2 Random:

Player 1 Win Rate: 63.50%

Player 2 Win Rate: 36.50%

P1 Random vs P2 Strategic:

Player 1 Win Rate: 26.30%


Player 2 Win Rate: 73.70%

Both Random:

Player 1 Win Rate: 48.10%

Player 2 Win Rate: 51.90%

Starting position: abidjan
 Move 0 (Strategic): ningbo - 21 moves available
 Move 1 (Random): ottawa - 5 moves available
 Move 2 (Strategic): adelaide - 33 moves available
 Move 3 (Random): esfahan - 3 moves available
 Move 4 (Strategic): nairobi - 20 moves available
 Move 5 (Random): istanbul - 6 moves available
 Move 6 (Strategic): lahore - 26 moves available
 Move 7 (Random): edmonton - 2 moves available
 Move 8 (Strategic): new taipei - 19 moves available
 Move 9 (Random): ibadan - 5 moves available
 Move 10 (Strategic): nnewi - 18 moves available
 Move 11 (Random): islamabad - 4 moves available
 Move 12 (Strategic): daegu - 22 moves available
 Move 13 (Random): uyo - 3 moves available
 Move 14 (Strategic): ouagadougou - 4 moves available
 Move 15 (Random): urumqi - 2 moves available
 Move 16 (Strategic): indore - 3 moves available
 Move 17 (Random): ekurhuleni - 1 moves available
 Move 18 (Strategic): izmir - 2 moves available
 Move 19 (Random): raipur - 12 moves available
 Move 20 (Strategic): rome - 11 moves available
 Game Over at move 21 - no unvisited nodes available from rome



2. Entropy Minimization and Forced Predictability

Goal: Make moves that keep your options open while forcing your opponent into **predictable** and **limited** future choices.

Failed Sub-Strategy: Just Minimizing Opponent's Out-Degree

Idea:

Always pick moves that leave the opponent with as few responses as possible.

Why It Failed:

- Some low-out-degree nodes lead to high-out-degree successors.
- Fails in early/mid-game, when too many moves are available.
- Doesn't account for move popularity (players tend to pick familiar names).

Improvement: Shannon Entropy for Opponent Control

- Uses **Shannon entropy** to measure opponent unpredictability.
- Accounts for **popularity bias**, ensuring that even if the opponent has moves, **they are predictable**.


Simulation Results:

Both Entropy:
Player 1 Win Rate: 76.00%
Player 2 Win Rate: 24.00%

P1 Entropy vs P2 Random:
Player 1 Win Rate: 89.00%
Player 2 Win Rate: 11.00%

P1 Random vs P2 Entropy:
Player 1 Win Rate: 21.00%
Player 2 Win Rate: 79.00%

Starting position: qinhuangdao
Move 0 (Entropy): ouagadougou - 5 moves available
Move 1 (Random): ulaanbaatar - 3 moves available
Move 2 (Entropy): rio de janeiro - 12 moves available
Move 3 (Random): onitsha - 4 moves available
Move 4 (Entropy): antananarivo - 34 moves available
Move 5 (Random): omsk - 3 moves available
Move 6 (Entropy): kano - 29 moves available
Move 7 (Random): ottawa - 2 moves available
Move 8 (Entropy): aleppo - 33 moves available
Move 9 (Random): osaka - 1 moves available
Move 10 (Entropy): adelaide - 32 moves available
Move 11 (Random): esfahan - 3 moves available
Move 12 (Entropy): ningbo - 21 moves available
Game Over at move 13 - no unvisited nodes available from ningbo



3. Centralized Choke-Point Strategy

Goal: Guide the opponent toward a **pre-determined weak position** by forcing them through a strategically chosen bottleneck node.

Failed Sub-Strategy: Just Picking High Betweenness Nodes

Idea:

Always move to a **high betweenness** node, assuming it will give us control.

Why It Failed:

- Not all high-betweenness nodes are choke points.
- Fails to direct the opponent to a weak position.
- Doesn't consider game trajectory.

Improvement: Choke-Point Frequency and Composite Score

- Identifies **specific choke nodes** that appear in **most paths** toward a weak position.
 - Quantifies **choke effectiveness** using:
 - **Path frequency** (how often the node appears in critical paths).
 - **Betweenness centrality** (ensuring structural importance).
 - **Out-degree filtering** (to prevent self-traps).
-

4. Local Lookahead Strategy

Goal: Simulate several turns in advance to find moves that give a **long-term advantage** while limiting the opponent's flexibility.

Failed Sub-Strategy: Greedy Move Maximization

Idea:

Always pick the move that maximizes your future options immediately.

Why It Failed:

- Doesn't consider opponent responses.

- **Fails in long-term strategy:** Some moves appear strong initially but lead to weak positions later.
- **Ignores network structure changes.**

Improvement: Multi-Metric Lookahead

- **Simulates 2–3 turns ahead** to evaluate move impact.
- **Evaluates:**
 - **Opponent's Future Options**
 - **Our Future Options**
 - **Spectral Gap Changes** (preventing self-trapping).

Simulation Results:

P1 Lookahead vs P2 Random:

Player 1 Win Rate: 66.67%

Player 2 Win Rate: 33.33%

P1 Random vs P2 Lookahead:

Player 1 Win Rate: 33.33%

Player 2 Win Rate: 66.67%

Candidate: iceland, Score: 0.3017, Metrics: {'my_spec_gap': 0.6737109800556802, 'avg_opp_spec': 3.2, 'delta_spec': -2.52628901994432, 'avg_my_moves': 10.944010695187165, 'avg_opp_moves': 9.4, 'delta_moves': 1.5440106951871648, 'H_self': 3.2679237614023524, 'H_opp': 1.9839779803810733, 'delta_entropy': 1.283945781021279, 'composite_score': 0.3016674562641237}

Candidate: india, Score: -6.9200, Metrics: {'my_spec_gap': 0.8767876410218791, 'avg_opp_spec': 7.545454545454543, 'delta_spec': -6.668666904432664, 'avg_my_moves': 11.097520661157025, 'avg_opp_moves': 11.181818181818182, 'delta_moves': -0.08429752066115626, 'H_self': 3.2914217751774415, 'H_opp': 3.458417206433648, 'delta_entropy': -0.1669954312562063, 'composite_score': -6.919959856350026}

Candidate: indonesia, Score: -6.9200, Metrics: {'my_spec_gap': 0.8767876410218791, 'avg_opp_spec': 7.545454545454543, 'delta_spec': -6.668666904432664, 'avg_my_moves': 11.097520661157025, 'avg_opp_moves': 11.181818181818182, 'delta_moves': -0.08429752066115626, 'H_self': 3.2914217751774415, 'H_opp': 3.458417206433648, 'delta_entropy': -0.1669954312562063, 'composite_score': -6.919959856350026}

Candidate: iran, Score: -3.3591, Metrics: {'my_spec_gap': 0.8810141427021404, 'avg_opp_spec': 5.545454545454546, 'delta_spec': -4.664440402752406, 'avg_my_moves': 11.650032795487341, 'avg_opp_moves': 10.636363636363637, 'delta_moves': 1.0136691591237046, 'H_self': 3.4627807864601214, 'H_opp': 3.1710760287075006, 'delta_entropy': 0.29170475775262084, 'composite_score': -3.35906648587608}

Candidate: iraq, Score: 11.5850, Metrics: {'my_spec_gap': 3.0000000000000004, 'avg_opp_spec': 0.9999999999999998, 'delta_spec': 2.0000000000000001, 'avg_my_moves': 12.0, 'avg_opp_moves': 4.0, 'delta_moves': 8.0, 'H_self': 1.584962500721156, 'H_opp': 0.0, 'delta_entropy': 1.584962500721156, 'composite_score': 11.584962500721156}

Candidate: ireland, Score: 0.3017, Metrics: {'my_spec_gap': 0.6737109800556802, 'avg_opp_spec': 3.2, 'delta_spec': -2.52628901994432, 'avg_my_moves': 10.944010695187165, 'avg_opp_moves': 9.4, 'delta_moves': 1.5440106951871648, 'H_self': 3.2679237614023524, 'H_opp': 1.9839779803810733, 'delta_entropy': 1.283945781021279, 'composite_score': 0.3016674562641237}

Candidate: israel, Score: -4.6916, Metrics: {'my_spec_gap': 1.4893305405599224, 'avg_opp_spec': 5.555555555555555, 'delta_spec': -4.0662250149956325, 'avg_my_moves': 11.272727272727273, 'avg_opp_moves': 12.555555555555555, 'delta_moves': -1.2828282828282828, 'H_self': 3.621462645940613, 'H_opp': 2.964019840191667, 'delta_entropy': 0.6574428057489463, 'composite_score': -4.691610492074968}

Candidate: italy, Score: 1.6842, Metrics: {'my_spec_gap': 0.8767876410218791, 'avg_opp_spec': 0.9999999999999993, 'delta_spec': -0.12321235897812022, 'avg_my_moves': 10.636363636363637, 'avg_opp_moves': 12.0, 'delta_moves': -1.3636363636363633, 'H_self': 3.1710760287075006, 'H_opp': 0.0, 'delta_entropy': 3.1710760287075006, 'composite_score': 1.684227306093017}

Best move from haiti: iraq with composite score 11.5850



Task 2 - Community Detection

Infomap

1. Modeling Game Flow as Information Flow

- **Random Walker Analogy**

- Infomap treats the game as a flow of information by simulating a random walker that moves from one country to another following the atlas rule. For instance, if the walker

is at “India” (ending with “a”), it can only move to countries that begin with “A” (like “Australia”).

- **Direction Matters**

- Since the atlas graph is directed, Infomap respects the game’s rule: moves only happen in the allowed direction (last letter → first letter). This ensures that the flow modeled by the algorithm exactly mimics the possible moves in the game.
-

2. The Map Equation and Community Detection

- **Minimizing Description Length** Infomap uses an information-theoretic concept called the map equation. It seeks to partition the graph into communities (clusters of countries) in such a way that the description length of a random walker’s journey is minimized. If the walker tends to stay within a certain group of countries, that group is identified as a community.
 - **Two-Level Coding** The algorithm assigns codes at two levels:
 - **Module (Community) Level:** Each community gets a unique code.
 - **Node Level:** Within each community, each country is given its own code. When the walker moves within a community, only the node-level code is needed. However, moving between communities requires switching to a module-level code. A good community structure is one where most moves are internal, leading to an efficient (compressed) description.
-

1. Interpreting the Computed Communities

- **Natural Groupings:** In the atlas graph, communities detected by Infomap often represent clusters of countries that are closely interlinked by the game’s rules. For example:
 - **Letter Patterns:**
 - Countries might group together because many have names that start or end with similar letters. A community could consist mostly of countries that start with “A” or end with “n,” reflecting the natural bias in letter transitions.
 - **Choke Points:**
 - Some countries—like Yemen in our example—can act as bridges. Yemen, with incoming edges from countries ending in “y” and outgoing edges to those starting with “n,” might form a critical part of a community or even lie at the interface between communities. Such nodes have high betweenness and become strategic choke points.
- **Alignment with Human Intuition:**

- When you look at the computed communities, you might notice that they correspond to groupings that humans would naturally recognize. People often group items based on common features; here, countries with similar starting or ending letters, or those that are strategically central in the flow of the game, naturally fall into the same clusters.
-

2. Strategic Implications for the Game of Atlas

- **Within-Community Moves:**
 - If you know that a group of countries forms a tightly connected community, playing within that community can be a safe strategy. It means you have many follow-up moves because the transitions (edges) are abundant within the community. This strategy minimizes the risk of leaving your opponent with an immediate winning move.
- **Exploiting Choke Points:** The Infomap algorithm might reveal that certain countries serve as bridges between communities. Choosing a country that lies on the border (or between) communities could force your opponent into a corner:
 - **Limiting Options:**
 - By steering the game through a choke point, you effectively funnel the game into a region where your opponent's valid moves are more restricted.
 - **Strategic Advantage:**
 - Even if you cannot immediately play the “ideal” country (for example, you may want to say “Hungary” but the current flow does not allow it), you can opt for a move that guides the game toward that region. This long-term planning—based on the community structure—gives you a competitive edge.
- **Intuition of a “Good” Strategy:**
 - The communities provide insight into the structure of valid moves:
 - **Control the Flow:**
 - By understanding which countries are central within a community or connect multiple communities, you can prioritize moves that maximize your control over the game's progression.
 - **Predictability:**
 - A move within a dense community might give you more predictable follow-up moves, while a move that transitions between communities could be used to disrupt your opponent's expected path.

GAE with GCN and HDBSCAN

1. What is GAE (Graph Autoencoder)?

A **Graph Autoencoder (GAE)** is a **neural network model** that learns **low-dimensional representations (embeddings)** of nodes in a graph while preserving structural relationships.

- **Encoder:** Compresses node features into a **latent space**.
 - **Decoder:** Reconstructs edges from embeddings, ensuring learned representations retain graph structure.
 - **Loss Function:** Minimizes reconstruction error, ensuring embeddings capture key node relationships.
-

2. Using GCN (Graph Convolutional Networks) as the Encoder

A **GCN (Graph Convolutional Network)** is a **type of neural network designed for graphs**, which learns node representations by aggregating information from neighbors.

In **GAE for the Game of Atlas**:

- **Nodes:** Countries
- **Edges:** Valid game transitions (last-letter to first-letter rule).
- **Node Features:** A **52-dimensional one-hot encoding** (first and last letter encoded separately).

GCN Encoder Layers:

1. **First Layer:** Aggregates neighboring features to learn local structure.
 2. **Activation (ReLU):** Introduces non-linearity.
 3. **Second Layer:** Maps to a **latent space (4D vector per country)**.
-

3. HDBSCAN: Clustering the Learned Embeddings

Once the **GCN-based encoder** learns **low-dimensional embeddings** for each country, we apply **HDBSCAN (Hierarchical Density-Based Clustering)** to group similar countries.
